

Single Dell Data Analysis Course

Data preprocessing

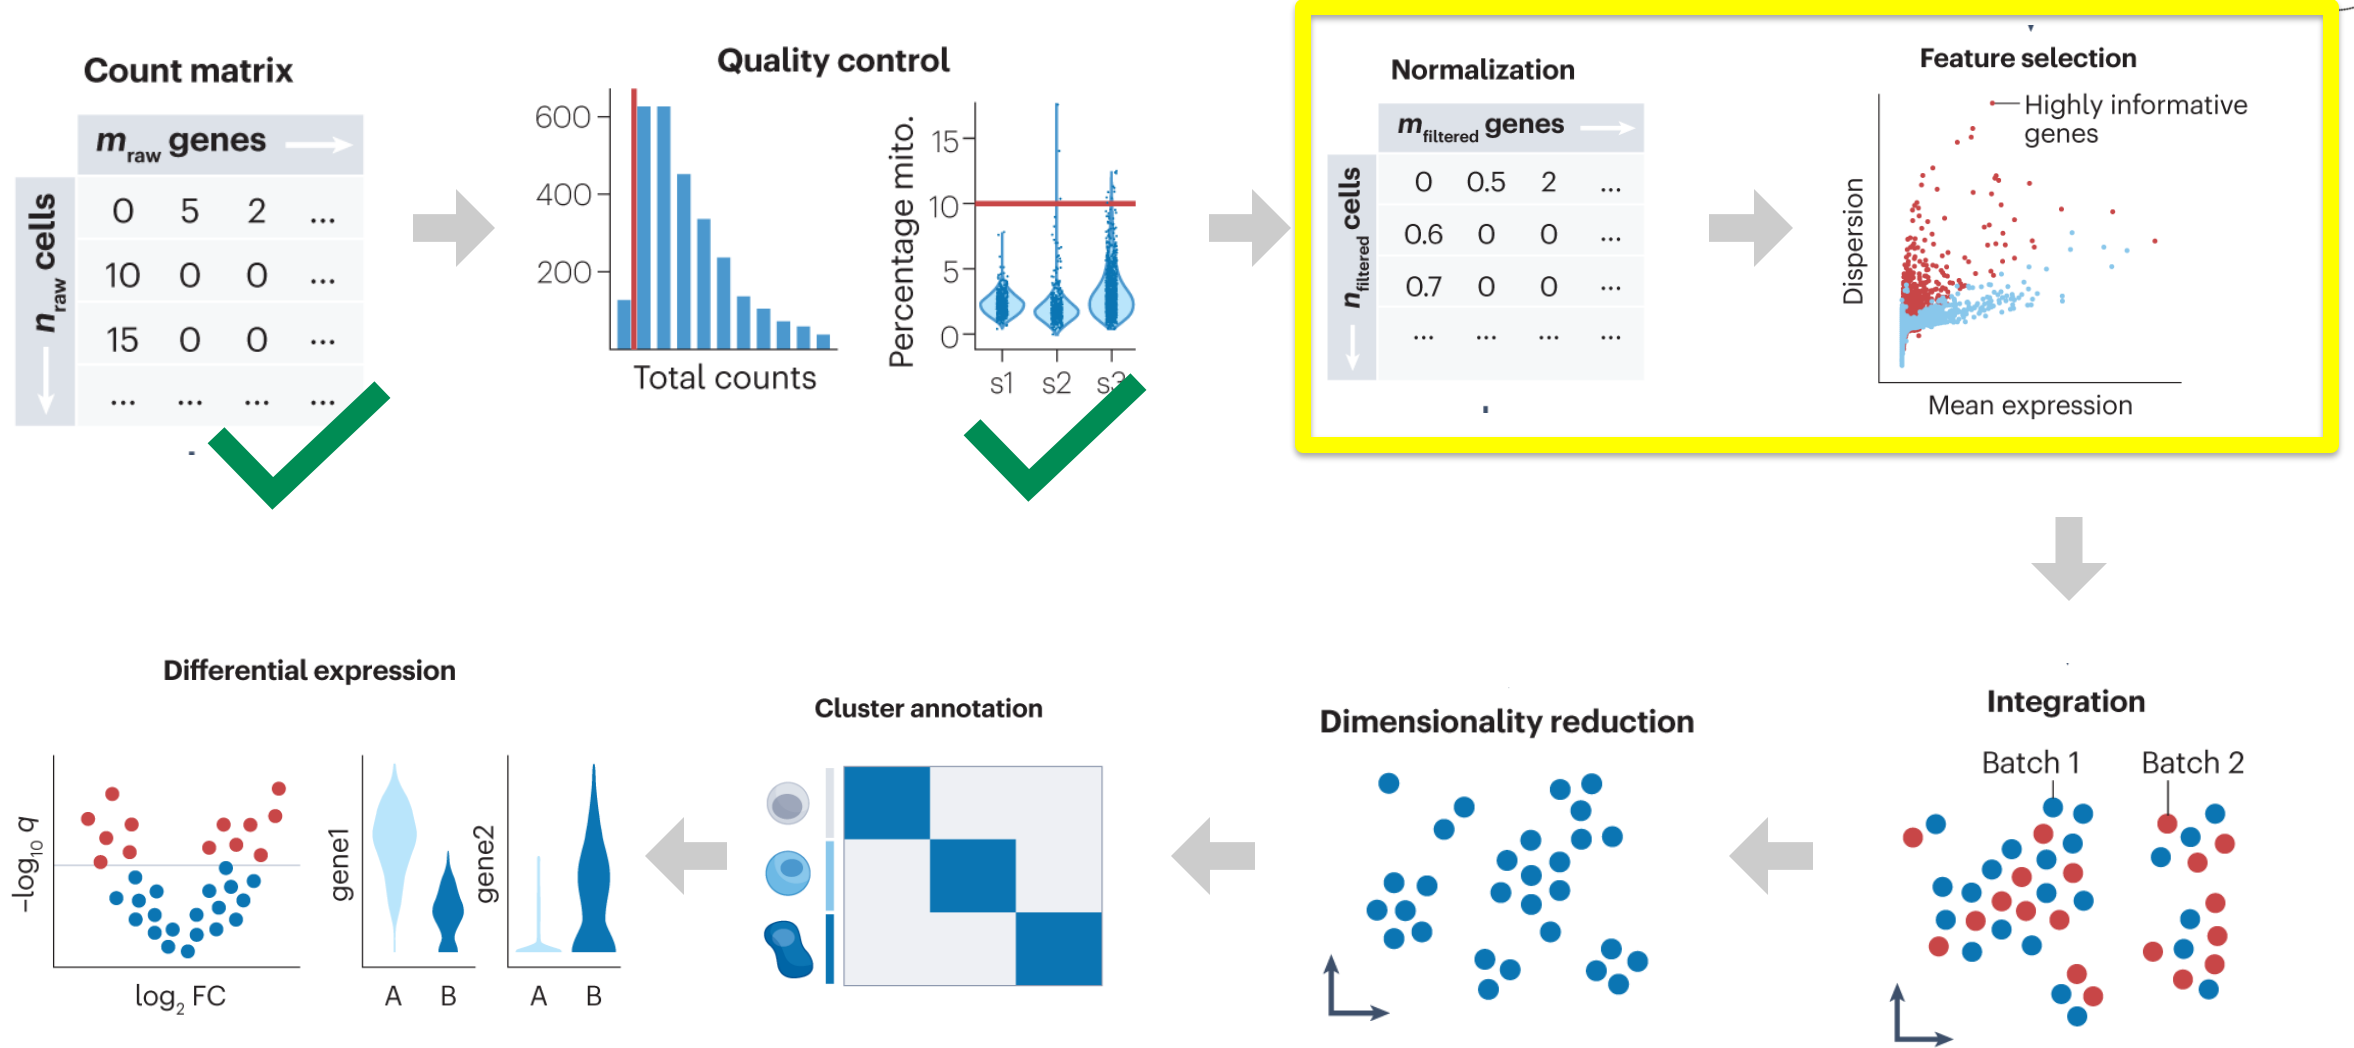
Lisa Buchauer

Professor of Systems Biology of Infectious Diseases

Department of Infectious Diseases and Intensive Care

Charité - Universitätsmedizin Berlin

Processing overview



Heumos, L., Schaar, A.C., Lance, C. et al. Best practices for single-cell analysis across modalities. Nat Rev Genet 24, 550–572 (2023). <https://doi.org/10.1038/s41576-023-00586-w>



Three important lines of code

Total-count normalize (library-size correct) the data matrix \mathbf{X} to 10,000 reads per cell, so that counts become comparable among cells.

1

```
sc.pp.normalize_total(adata, target_sum=1e4)
```

```
normalizing counts per cell  
finished (0:00:00)
```

Logarithmize the data:

2

```
sc.pp.log1p(adata)
```

Identify highly-variable genes.

3

```
sc.pp.highly_variable_genes(adata, min_mean=0.0125, max_mean=3, min_disp=0.5)
```

Normalizing data to mitigate library size effects



scanpy

```
sc.pp.normalize_total(adata, target_sum=1e4)
```

Seurat

```
pbmc <- NormalizeData(pbmc, normalization.method = "LogNormalize", scale.factor = 10000)
```



Normalizing data to mitigate library size effects

scanpy




```
sc.pp.normalize_total(adata, target_sum=1e4)
```

Seurat




```
pbmc <- NormalizeData(pbmc, normalization.method = "LogNormalize", scale.factor = 10000)
```

Why?

in the cells



in the count
matrix



“drop-out”

Normalizing data to mitigate library size effects



How?

target_sum = 10

(count / library
size) x target sum

A|G1: $(\frac{1}{2}) * 10 = 5$

B|G2: $(\frac{4}{8}) * 10 = 5$

	A	B	C	D	E
G1	1	4	0	1	4
G2	1	4	2	3	2
G3	0	0	4	3	2
Library Size (Σ)	2	8	6	7	8

	A	B	C	D	E
G1	5	5	0	1.43	5
G2	5	5	3.33	4.29	2.5
G3	0	0	6.67	4.29	2.5
Sum	10	10	10	10	10

Normalizing data to mitigate library size effects



How?

	A	B	C	D	E
G1	1	4	0	1	4
G2	1	4	2	3	2
G3	0	0	4	3	2
Library Size (Σ)	2	8	6	7	8

target_sum = 10

(count / library size) x target sum

A|G1: $(1/2) * 10 = 5$

B|G2: $(4/8) * 10 = 5$

	A	B	C	D	E
G1	5	5	0	1.43	5
G2	5	5	3.33	4.29	2.5
G3	0	0	6.67	4.29	2.5
Sum	10	10	10	10	10

- 10k counts are often used as target
- Alternative: normalize to median library size of original data

- Relies on the assumption that every cell originally had the same amount of RNA
- ! Actual variation in count number may be due to both technical AND biological effects



Taking the log to stabilize the variance

scanpy

```
sc.pp.log1p(adata)
```

Seurat

```
pbmc <- NormalizeData(pbmc, normalization.method = "LogNormalize", scale.factor = 10000)
```




scanpy



```
sc.pp.log1p(adata)
```

Seurat

```
pbmc <- NormalizeData(pbmc, normalization.method = "LogNormalize", scale.factor = 10000)
```

Why?

Many downstream methods like identification of highly variable genes, dimension reduction and clustering require (or at least perform a lot better) with **homoscedastic** data.





Taking the log to stabilize the variance

scanpy



```
sc.pp.log1p(adata)
```

Seurat

```
pbmc <- NormalizeData(pbmc, normalization.method = "LogNormalize", scale.factor = 10000)
```

Why?


Many downstream methods like identification of highly variable genes, dimension reduction and clustering expect (or at least perform a lot better with) **homoscedastic** data.





Why?


Taking the log to stabilize the variance



Lowly expressed gene

Mean = 1 count

Variance = 0.4 counts



Highly expressed gene


Mean = 11.4 count

Variance = 6.64 counts



Why?

Taking the log to stabilize the variance



Lowly expressed gene

Mean = 1 count


Variance = 0.4 counts

Highly expressed gene

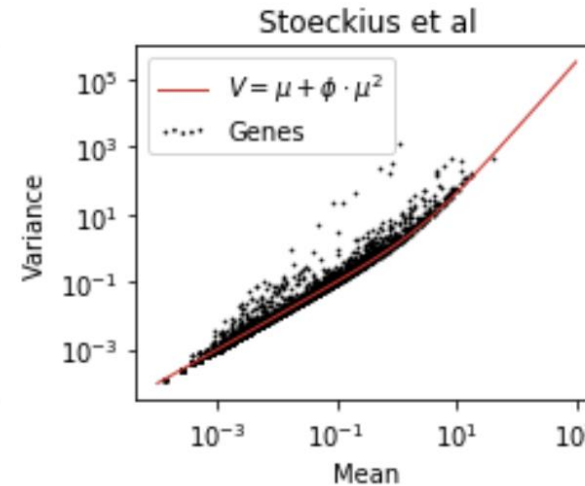


Mean = 11.4 count

Variance = 6.64 counts

homoscedasticity



heteroscedasticity

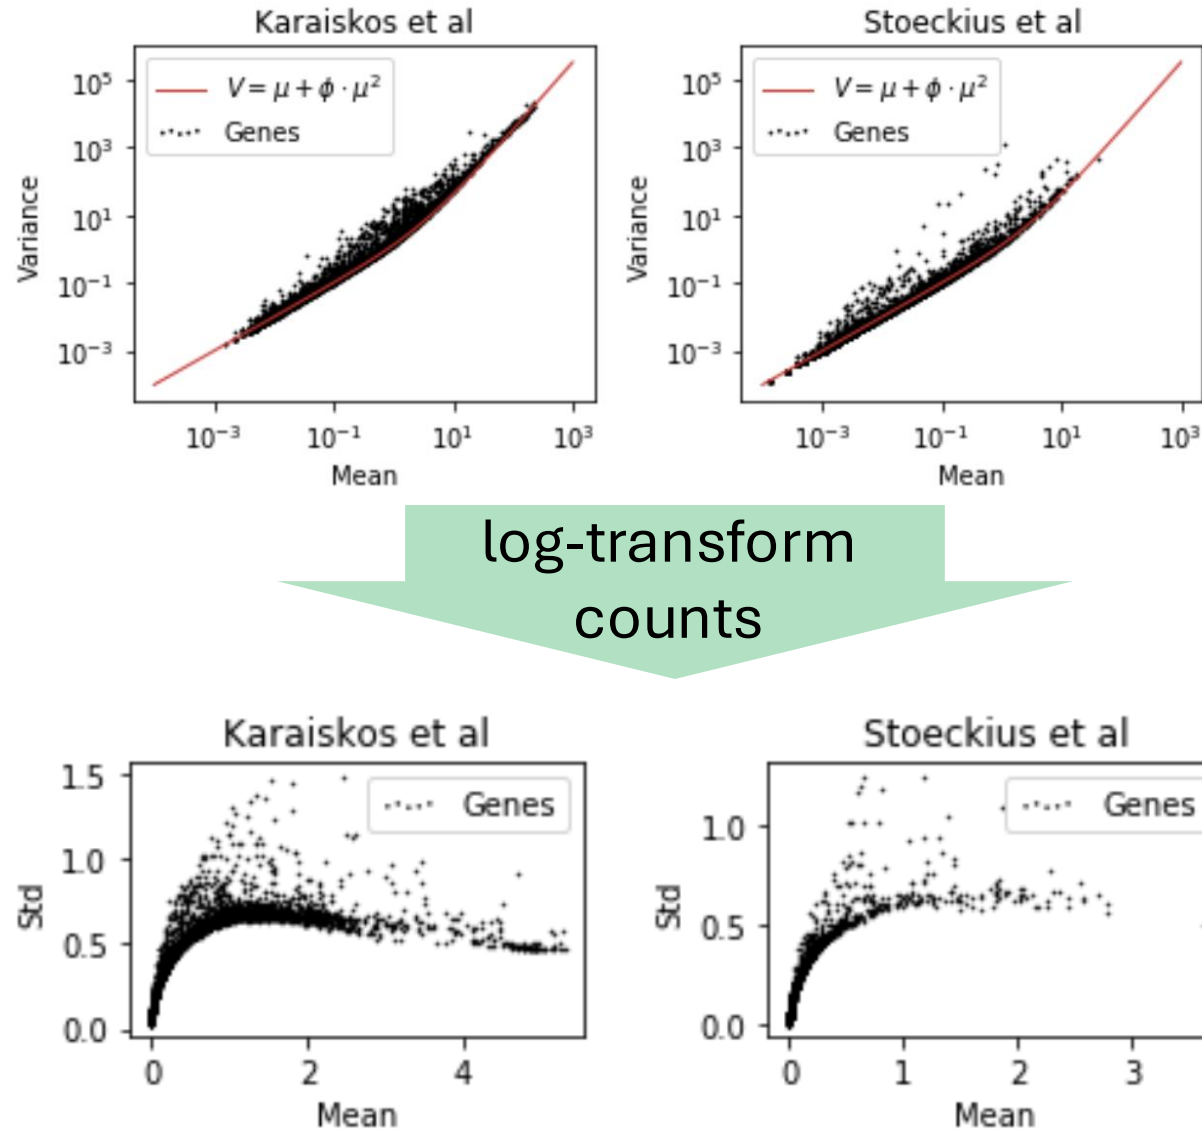


Heteroscedastic!




Taking the log to stabilize the variance


Why?



homoscedasticity



heteroscedasticity



Taking the log to stabilize the variance



How?

	A	B	C	D	E
G1	5	5	0	1.43	5
G2	5	5	3.33	4.29	2.5
G3	0	0	6.67	4.29	2.5
Sum	10	10	10	10	10

log(counts)

Taking the log to stabilize the variance



How?

	A	B	C	D	E
G1	5	5	0	1.43	5
G2	5	5	3.33	4.29	2.5
G3	0	0	6.67	4.29	2.5
Sum	10	10	10	10	10

log(counts)

Value Error

Taking the log to stabilize the variance




How?

	A	B	C	D	E
G1	5	5	0	1.43	5
G2	5	5	3.33	4.29	2.5
G3	0	0	6.67	4.29	2.5
Sum	10	10	10	10	10

log(counts)

Value Error



Taking the log to stabilize the variance



How?

	A	B	C	D	E
G1	5	5	0	1.43	5
G2	5	5	3.33	4.29	2.5
G3	0	0	6.67	4.29	2.5
Sum	10	10	10	10	10

$\log(\text{counts} + 1)$

	A	B	C	D	E
G1	1.8	1.8	0	0.9	1.8
G2	1.8	1.8	1.5	1.7	0.9
G3	0	0	2.0	1.7	0.9

Taking the log to stabilize the variance



How?

	A	B	C	D	E
G1	5	5	0	1.43	5
G2	5	5	3.33	4.29	2.5
G3	0	0	6.67	4.29	2.5
Sum	10	10	10	10	10

$\log(\text{counts} + 1)$

	A	B	C	D	E
G1	1.8	1.8	0	0.9	1.8
G2	1.8	1.8	1.5	1.7	0.9
G3	0	0	1.7	1.7	0.9

Lognormalized data,
ready for downstream
processing



- Logarithmic transformation is the most common choice for this task
- Many alternatives exist, but performance differences are minor

Finding highly variable genes



scanpy

```
sc.pp.highly_variable_genes(adata, n_top_genes=2000, batch_key="sample")
```

Seurat

```
pbmc <- FindVariableFeatures(pbmc, selection.method = "vst", nfeatures = 2000)
```



Finding highly variable genes

scanpy

```
sc.pp.highly_variable_genes(adata, n_top_genes=2000, batch_key="sample")
```

Seurat

```
pbmc <- FindVariableFeatures(pbmc, selection.method = "vst", nfeatures = 2000)
```

Why?

	A	B	C	D	E
G4	0	0	0	0	0
G5	1.8	1.9	0.2	0.2	2.1
G6	0.5	0.5	1.0	0.9	0.4
G7	0	0	0.1	0	0
G8	1.6	1.5	1.6	1.5	1.7



Finding highly variable genes

scanpy

```
sc.pp.highly_variable_genes(adata, n_top_genes=2000, batch_key="sample")
```

Seurat

```
pbmc <- FindVariableFeatures(pbmc, selection.method = "vst", nfeatures = 2000)
```

Why?


	A	B	C	D	E
G4	0	0	0	0	0
G5	1.8	1.9	0.2	0.2	2.1
G6	0.5	0.5	1.0	0.9	0.4
G7	0	0	0.1	0	0
G8	1.6	1.5	1.6	1.5	1.7

- Genes that are hardly expressed at all and/or do not vary a lot across cells are less valuable for analysis
- Masking them increases computational efficiency while simultaneously reducing analysis noise



Finding highly variable genes

How?




Highly expressed gene

Mean = 11.4 count

Variance = 6.64 counts


Naïve idea: Let's just take the genes with the highest variance across cells.





How?

Finding highly variable genes




Highly expressed gene

Mean = 11.4 count


Variance = 6.64 counts

Naïve idea: Let's just take the genes with the highest variance across cells.



Heteroscedastic!

- Genes with higher expression also have higher variance by default
- To find the interesting genes, we need to compare their variability with that of similarly expressed genes



Finding highly variable genes



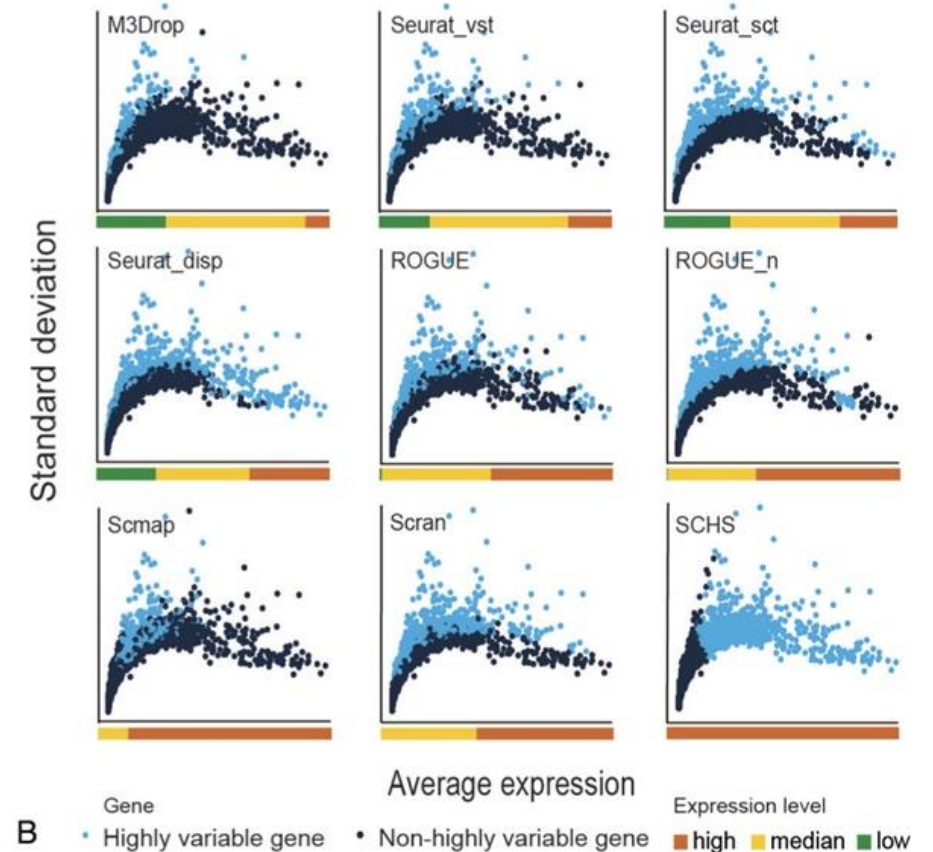
How?

typical strategy

Stabilize variance
(e.g. log-normalize counts)

Divide genes into bins based on
expression **or** fit a curve to the
standard deviation over mean
expression

Select genes which have higher
variance than their peers



Zhang, Yanan; Xie, Xiaowei; Wu, Peng*; Zhu, Ping*. SIEVE: identifying robust single cell variable genes for single-cell RNA sequencing data. Blood Science 3(2):p 35-39, April 2021. | DOI: 10.1097/BS9.0000000000000072

4?

Scaling / z-scoring / standardization brings all gene expression values onto a common scale



All Seurat
tutorials:

```
all.genes <- rownames(pbmc)  
pbmc <- ScaleData(pbmc, features = all.genes)
```

Most scanpy
tutorials:

[]:

4?

Scaling / z-scoring / standardization brings all gene expression values onto a common scale



All Seurat tutorials:

```
all.genes <- rownames(pbmc)
pbmc <- ScaleData(pbmc, features = all.genes)
```

Most scanpy tutorials:

[]:

$$x = \frac{x - \text{mean}(\vec{x})}{\text{std}(x)}$$

- Subtract mean
- Divide by standard deviation
- → all values are centered around 0 and have unit variance

4?

Scaling / z-scoring / standardization brings all gene expression values onto a common scale

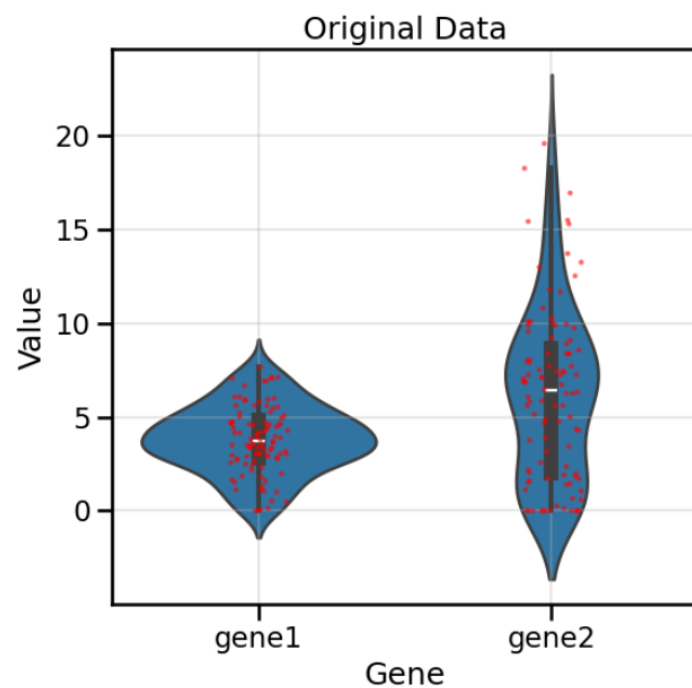


All Seurat tutorials:

```
all.genes <- rownames(pbmc)
pbmc <- ScaleData(pbmc, features = all.genes)
```

Most scanpy tutorials:

[]:



4?

Scaling / z-scoring / standardization brings all gene expression values onto a common scale




All Seurat tutorials:

```
all.genes <- rownames(pbmc)
pbmc <- ScaleData(pbmc, features = all.genes)
```

Most scanpy tutorials:

[]:



4?

Scaling / z-scoring / standardization brings all gene expression values onto a common scale




All Seurat tutorials:

```
all.genes <- rownames(pbmc)
pbmc <- ScaleData(pbmc, features = all.genes)
```

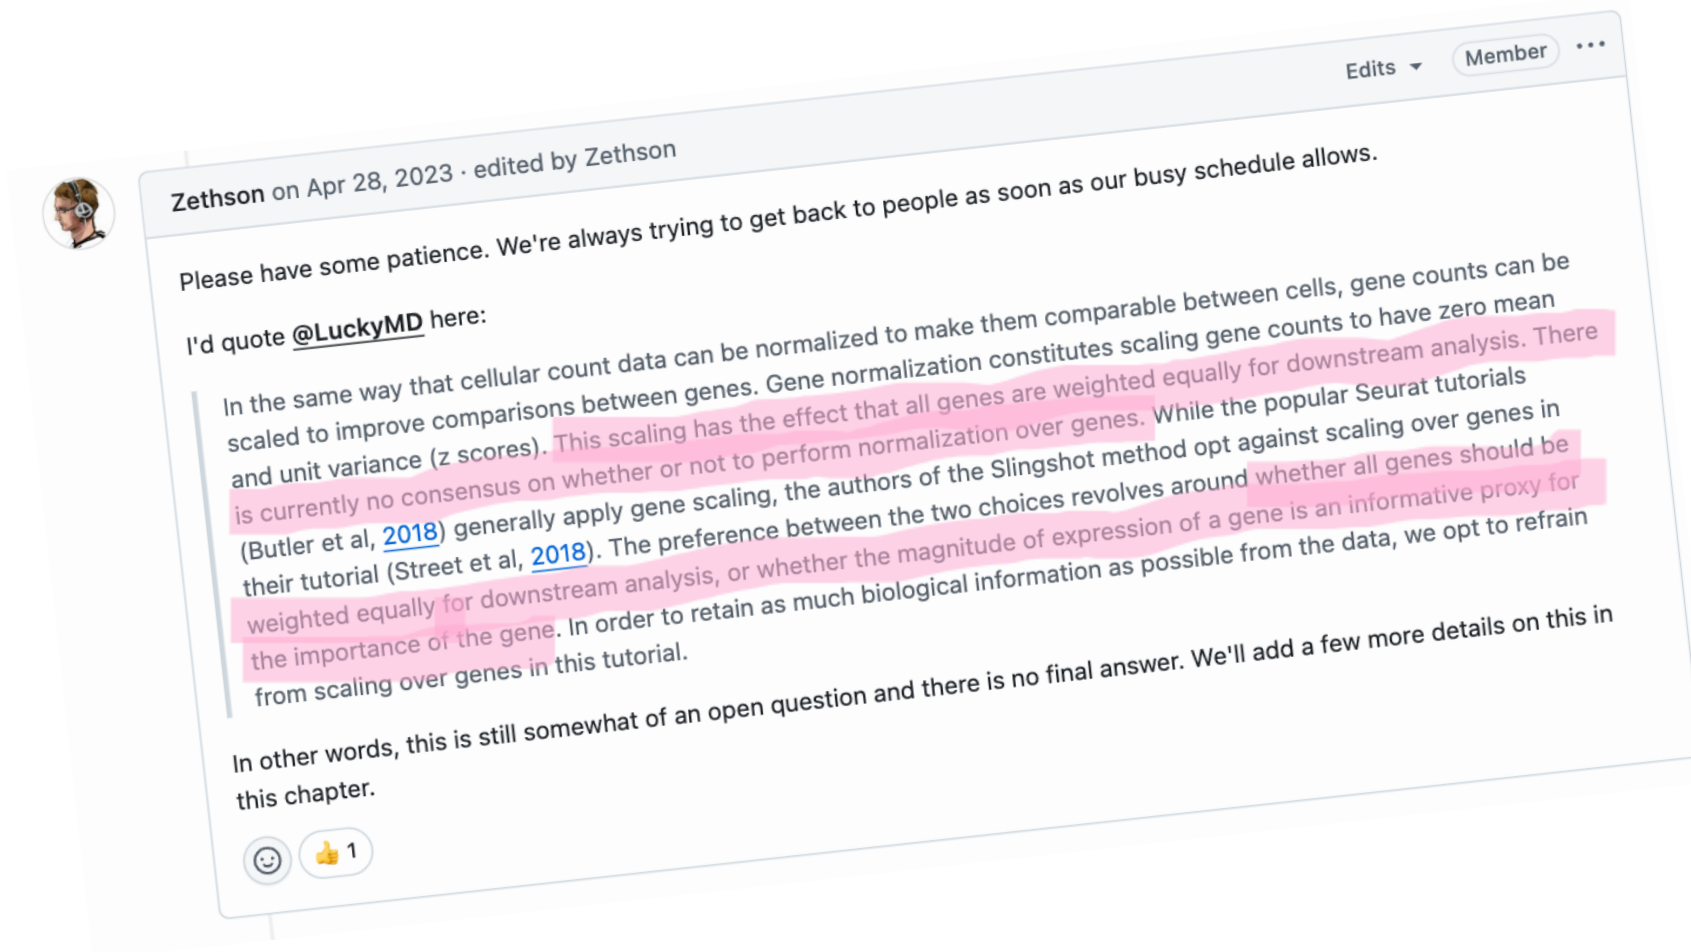
Most scanpy tutorials:

[]:





To scale or not to scale?





No!

- Preserves relative variance differences between genes
- Highly variable genes naturally dominate PCs
- More similar to traditional bulk RNA-seq PCA

Aye!

- All genes contribute equally (unit variance)
- Prevents highly expressed genes from dominating just due to magnitude
- Can reveal subtle patterns in lower-variance genes



No!

- Preserves relative variance differences between genes
- Highly variable genes naturally dominate PCs
- More similar to traditional bulk RNA-seq PCA

scanpy

Aye!

- All genes contribute equally (unit variance)
- Prevents highly expressed genes from dominating just due to magnitude
- Can reveal subtle patterns in lower-variance genes


Seurat

(you can hack Seurat to circumvent scaling, but this leads downstream complications because the functions there expect scaled data, so you would also need to handle those)



Example of effects upcoming in workshop data

Split
appears
only
without
scaling!



- Non-scaling approach reveals a “duplicate” population
- The difference between the twins is just one highly expressed gene
- Through scaling, the importance of this gene is lessened and the split does not appear on the UMAP



1

2

3

4

Alternative: Pearson Residuals (SCTransform)

scanpy

```
sc.experimental.pp.recipe_pearson_residuals(adata)
```

Seurat

```
# run sctransform  
pbmc <- SCTransform(pbmc)
```

Replace normalization, log-transformation and highly variable gene search

Why?

Total gene
expression
variability

=

Technical
variability

+

Biological
variability

https://scanpy-tutorials.readthedocs.io/en/latest/tutorial_pearson_residuals.html
https://satijalab.org/seurat/articles/sctransform_vignette



1

2

3

4

Alternative: Pearson Residuals (SCTransform)

scanpy

```
sc.experimental.pp.recipe_pearson_residuals(adata)
```

Seurat

```
# run sctransform  
pbmc <- SCTransform(pbmc)
```

Replace normalization, log-transformation and highly variable gene search

Why?

Total gene
expression
variability

=

Technical
variability

+

Biological
variability

Biological
variability

=

Total gene
expression
variability

-

Technical
variability

measure

stats model

https://scanpy-tutorials.readthedocs.io/en/latest/tutorial_pearson_residuals.html
https://satijalab.org/seurat/articles/sctransform_vignette

1

2

3

4

Alternative: Pearson Residuals (SCTransform)



How?

Biological
variability

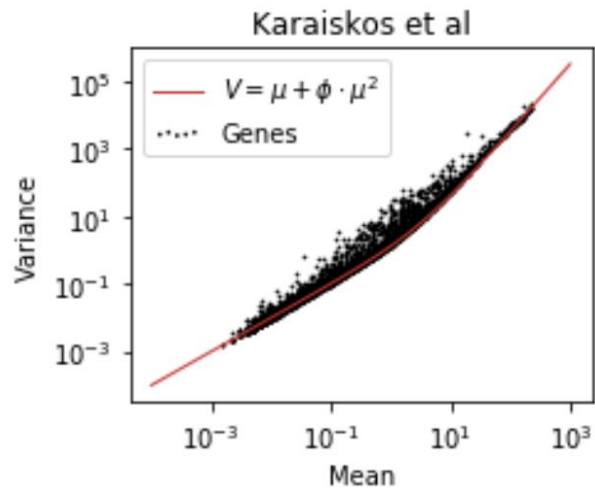
=

Total gene
expression
variability

measure

Technical
variability

stats model



Negative binomial
distribution describes the
technical noise of single
cell data.

Pearson
residual

Raw
count

Expected
count under
NB model

$$Z_{cg} = \frac{X_{cg} - \hat{\mu}_{cg}}{\sqrt{\hat{\mu}_{cg} + \hat{\mu}_{cg}^2 / \theta}}$$

Expected std
under NB model

Lause, J., Berens, P. & Kobak, D. Analytic Pearson residuals for normalization of single-cell RNA-seq UMI data. Genome Biol 22, 258 (2021). <https://doi.org/10.1186/s13059-021-02451-7>


1

2

3

4

Alternative: Pearson Residuals (SCTransform)



→ after transformation into
Pearson Residuals

Basic Interpretation

- **Pearson residual = 0**: The observed count matches exactly what the model expected
- **Positive residual (> 0)**: The observed count is higher than expected
- **Negative residual (< 0)**: The observed count is lower than expected

Magnitude Interpretation

The magnitude of a Pearson residual indicates the strength of the deviation:

- **|residual| < 2** : Minor deviation, likely just random noise
- **|residual| > 3** : Strong deviation, highly likely to be biologically significant
- **|residual| > 5** : Extreme deviation, almost certainly represents a real biological signal




1

2

3

4

Alternative: Pearson Residuals (SCTransform)






1

2

3

4

Alternative: Pearson Residuals (SCTransform)



- Highly variable gene selection via Pearson Residuals can identify genes relevant for tiny cell populations
- Downstream analyses (e.g. clustering) can benefit from this
- Calculation is relatively expensive and can be prohibitive for large datasets

1

2

3

4

Pearson Residuals are not a must



Analysis | [Open access](#) | Published: 10 April 2023

Comparison of transformations for single-cell RNA-seq data

[Constantin Ahlmann-Eltze](#)  & [Wolfgang Huber](#)

[Nature Methods](#) **20**, 665–672 (2023) | [Cite this article](#)

38k Accesses | **10** Citations | **196** Altmetric | [Metrics](#)

Abstract

The count table, a numeric matrix of genes \times cells, is the basic input data structure in the analysis of single-cell RNA-sequencing data. A common preprocessing step is to adjust the counts for variable sampling efficiency and to transform them so that the variance is similar across the dynamic range. These steps are intended to make subsequent application of generic statistical methods more palatable. Here, we describe four transformation approaches based on the **delta method, model residuals**, inferred latent expression state and factor analysis. We compare their strengths and weaknesses and find that the latter three have appealing theoretical properties; however, in benchmarks using simulated and real-world data, it turns out that a rather simple approach, namely, **the logarithm with a pseudo-count followed by principal-component analysis, performs as well or better than the more sophisticated alternatives**. This result highlights limitations of current theoretical analysis as assessed by bottom-line performance benchmarks.

Ahlmann-Eltze, C., Huber, W. Comparison of transformations for single-cell RNA-seq data. *Nat Methods* 20, 665–672 (2023).
<https://doi.org/10.1038/s41592-023-01814-1>

To scale or not to scale?



Standard Workflows:

Scanpy (typical):

```
python  
sc.pp.normalize_total(adata) # normalize to counts per 10k  
sc.pp.log1p(adata) # log(1+x) transform  
sc.tl.pca(adata, zero_center=True) # PCA with centering  
No scaling to unit variance!
```

Seurat (typical):

```
r  
NormalizeData(obj) # similar log normalization  
ScaleData(obj) # centers AND scales to unit variance  
RunPCA(obj) # expects already-centered data
```

Your Situation:

You're feeding **log-normalized but unscaled** data to Seurat's PCA, which likely **doesn't center** because it assumes `ScaleData()` was already run. This would explain the difference!

To test your hypothesis:

1. Check if Seurat's `RunPCA()` has a centering parameter (I believe it doesn't by default)
2. You could manually center your data before feeding it to Seurat:
3. r
`data.centered <- sweep(data, 1, rowMeans(data), "-")` # Then run PCA on this
4. # After `NormalizeData`, manually center without scaling
5. Or check what scanpy's PCA looks like with `zero_center=False` - it should look more like your Seurat result

Why the Different Philosophies?

Scanpy approach (log-norm + center, no scaling):

- Preserves relative variance differences between genes
- Highly variable genes naturally dominate PCs
- More similar to traditional bulk RNA-seq PCA

Seurat approach (log-norm + scale):

- All genes contribute equally (unit variance)
- Prevents highly expressed genes from dominating just due to magnitude
- Can reveal subtle patterns in lower-variance genes

The difference in plots you're seeing is likely because:

- Scanpy: PCA on centered but unscaled log-counts
- Seurat (your hack): PCA on uncentered, unscaled log-counts

The centering makes a huge difference, as we discussed earlier - without it, PC1 often just captures mean expression level rather than biological variation.