

# Single Dell Data Analysis Course

## **Data preprocessing**

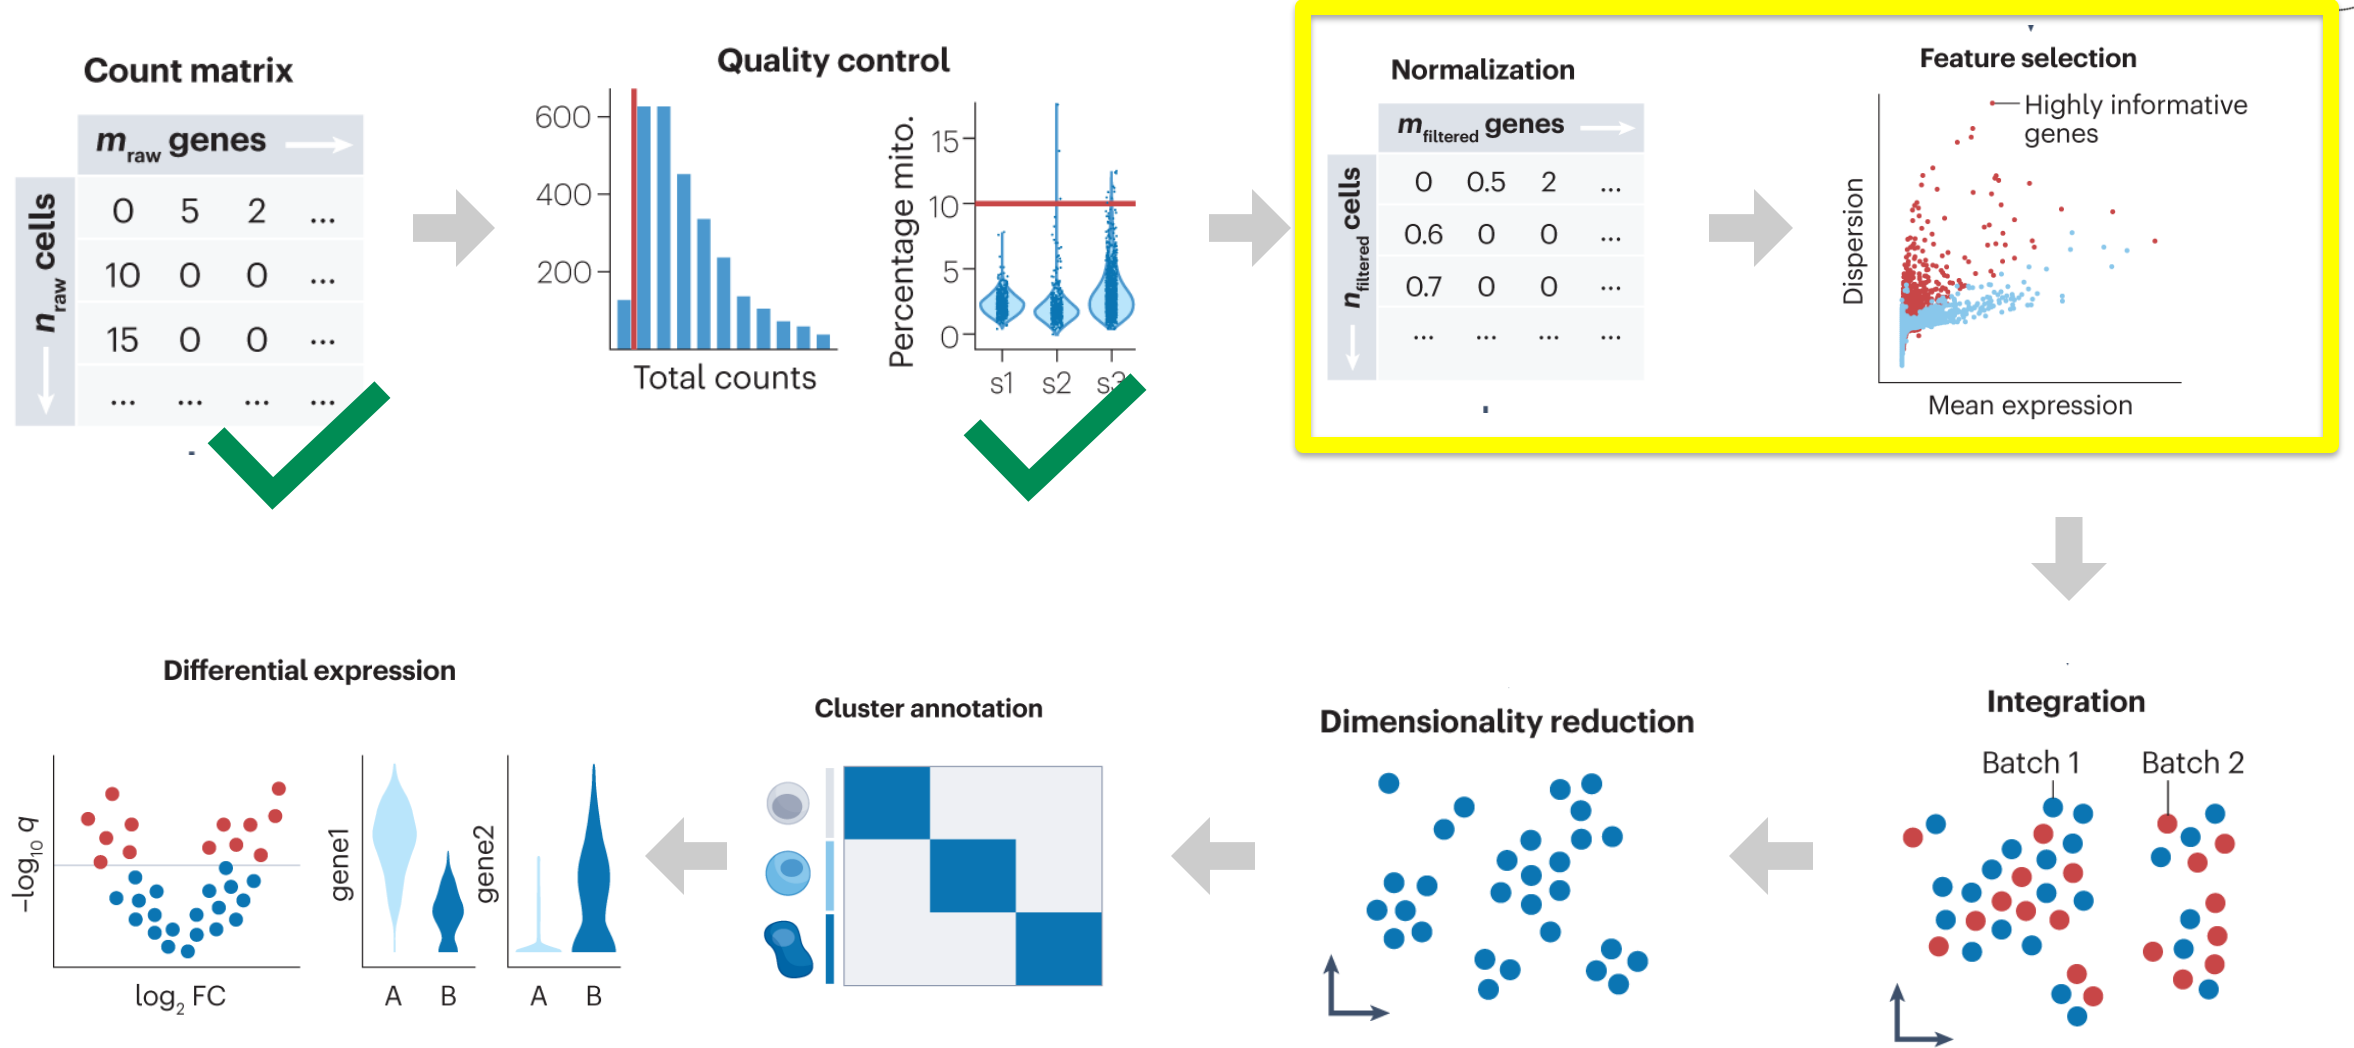
Lisa Buchauer

*Professor of Systems Biology of Infectious Diseases*

Department of Infectious Diseases and Intensive Care

Charité - Universitätsmedizin Berlin

# Processing overview



Heumos, L., Schaar, A.C., Lance, C. et al. Best practices for single-cell analysis across modalities. Nat Rev Genet 24, 550–572 (2023). <https://doi.org/10.1038/s41576-023-00586-w>



# Three important lines of code

Total-count normalize (library-size correct) the data matrix  $\mathbf{X}$  to 10,000 reads per cell, so that counts become comparable among cells.

1

```
sc.pp.normalize_total(adata, target_sum=1e4)
```

```
normalizing counts per cell  
finished (0:00:00)
```

Logarithmize the data:

2

```
sc.pp.log1p(adata)
```

Identify highly-variable genes.

3

```
sc.pp.highly_variable_genes(adata, min_mean=0.0125, max_mean=3, min_disp=0.5)
```

# Normalizing data to mitigate library size effects



scanpy

```
sc.pp.normalize_total(adata, target_sum=1e4)
```

Seurat

```
pbmc <- NormalizeData(pbmc, normalization.method = "LogNormalize", scale.factor = 10000)
```



# Normalizing data to mitigate library size effects

scanpy




```
sc.pp.normalize_total(adata, target_sum=1e4)
```

Seurat




```
pbmc <- NormalizeData(pbmc, normalization.method = "LogNormalize", scale.factor = 10000)
```

Why?

in the cells



in the count  
matrix



“drop-out”

# Normalizing data to mitigate library size effects



How?

target\_sum = 10

(count / library  
size) x target sum

A|G1:  $(\frac{1}{2}) * 10 = 5$

B|G2:  $(\frac{4}{8}) * 10 = 5$

	A	B	C	D	E
G1	1	4	0	1	4
G2	1	4	2	3	2
G3	0	0	4	3	2
Library Size ( $\Sigma$ )	2	8	6	7	8

	A	B	C	D	E
G1	5	5	0	1.43	5
G2	5	5	3.33	4.29	2.5
G3	0	0	6.67	4.29	2.5
Sum	10	10	10	10	10

# Normalizing data to mitigate library size effects



How?

	A	B	C	D	E
G1	1	4	0	1	4
G2	1	4	2	3	2
G3	0	0	4	3	2
Library Size ( $\Sigma$ )	2	8	6	7	8

target\_sum = 10

(count / library size) x target sum

A|G1:  $(1/2) * 10 = 5$

B|G2:  $(4/8) * 10 = 5$

	A	B	C	D	E
G1	5	5	0	1.43	5
G2	5	5	3.33	4.29	2.5
G3	0	0	6.67	4.29	2.5
Sum	10	10	10	10	10

- 10k counts are often used as target
- Alternative: normalize to median library size of original data

- Relies on the assumption that every cell originally had the same amount of RNA
- ! Actual variation in count number may be due to both technical AND biological effects



## Taking the log to stabilize the variance

scanpy

```
sc.pp.log1p(adata)
```

Seurat

```
pbmc <- NormalizeData(pbmc, normalization.method = "LogNormalize", scale.factor = 10000)
```





scanpy



```
sc.pp.log1p(adata)
```

Seurat

```
pbmc <- NormalizeData(pbmc, normalization.method = "LogNormalize", scale.factor = 10000)
```

Why?

Many downstream methods like identification of highly variable genes, dimension reduction and clustering require (or at least perform a lot better) with **homoscedastic** data.





## Taking the log to stabilize the variance

scanpy



```
sc.pp.log1p(adata)
```

Seurat

```
pbmc <- NormalizeData(pbmc, normalization.method = "LogNormalize", scale.factor = 10000)
```

Why?


Many downstream methods like identification of highly variable genes, dimension reduction and clustering expect (or at least perform a lot better with) **homoscedastic** data.





Why?


## Taking the log to stabilize the variance



**Lowly expressed gene**

Mean = 1 count

Variance = 0.4 counts



**Highly expressed gene**


Mean = 11.4 count

Variance = 6.64 counts



Why?

## Taking the log to stabilize the variance



### Lowly expressed gene

Mean = 1 count


Variance = 0.4 counts

### Highly expressed gene

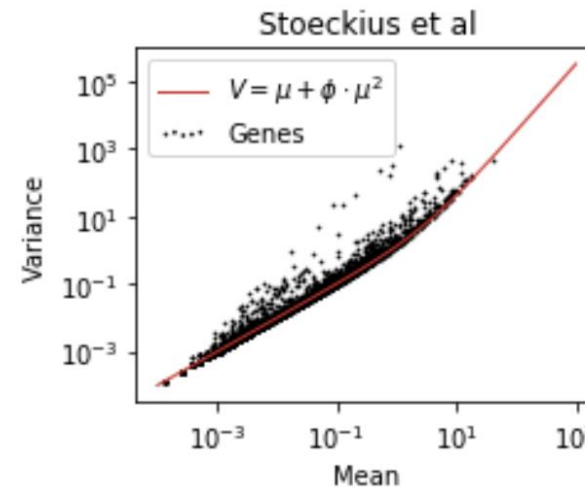


Mean = 11.4 count

Variance = 6.64 counts

homoscedasticity



heteroscedasticity

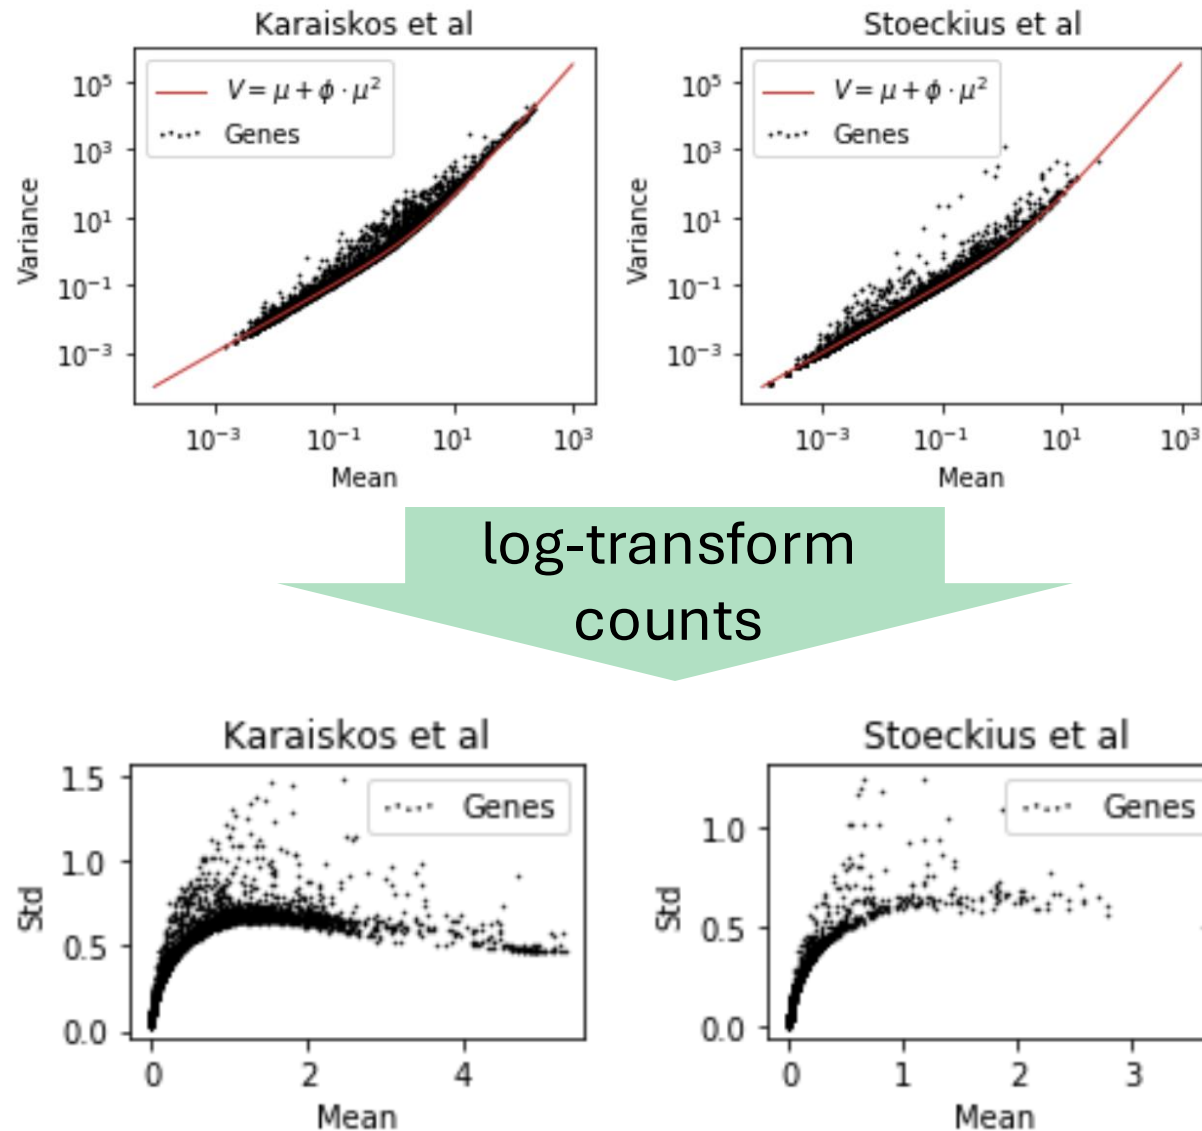


Heteroscedastic!




# Taking the log to stabilize the variance


Why?



homoscedasticity



heteroscedasticity



## Taking the log to stabilize the variance



How?

	A	B	C	D	E
G1	5	5	0	1.43	5
G2	5	5	3.33	4.29	2.5
G3	0	0	6.67	4.29	2.5
Sum	10	10	10	10	10

log(counts)

## Taking the log to stabilize the variance



How?

	A	B	C	D	E
G1	5	5	0	1.43	5
G2	5	5	3.33	4.29	2.5
G3	0	0	6.67	4.29	2.5
Sum	10	10	10	10	10

log(counts)

Value Error

# Taking the log to stabilize the variance




How?

	A	B	C	D	E
G1	5	5	0	1.43	5
G2	5	5	3.33	4.29	2.5
G3	0	0	6.67	4.29	2.5
Sum	10	10	10	10	10

log(counts)

Value Error





# Taking the log to stabilize the variance



How?

	A	B	C	D	E
G1	5	5	0	1.43	5
G2	5	5	3.33	4.29	2.5
G3	0	0	6.67	4.29	2.5
Sum	10	10	10	10	10

$\log(\text{counts} + 1)$

	A	B	C	D	E
G1	1.8	1.8	0	0.9	1.8
G2	1.8	1.8	1.5	1.7	0.9
G3	0	0	2.0	1.7	0.9

# Taking the log to stabilize the variance



How?

	A	B	C	D	E
G1	5	5	0	1.43	5
G2	5	5	3.33	4.29	2.5
G3	0	0	6.67	4.29	2.5
Sum	10	10	10	10	10

$\log(\text{counts} + 1)$

	A	B	C	D	E
G1	1.8	1.8	0	0.9	1.8
G2	1.8	1.8	1.5	1.7	0.9
G3	0	0	1.7	1.7	0.9

Lognormalized data,  
ready for downstream  
processing



- Logarithmic transformation is the most common choice for this task
- Many alternatives exist, but performance differences are minor

## Finding highly variable genes



scanpy

```
sc.pp.highly_variable_genes(adata, n_top_genes=2000, batch_key="sample")
```

Seurat

```
pbmc <- FindVariableFeatures(pbmc, selection.method = "vst", nfeatures = 2000)
```



## Finding highly variable genes

scanpy

```
sc.pp.highly_variable_genes(adata, n_top_genes=2000, batch_key="sample")
```

Seurat

```
pbmc <- FindVariableFeatures(pbmc, selection.method = "vst", nfeatures = 2000)
```

Why?

	A	B	C	D	E
G4	0	0	0	0	0
G5	1.8	1.9	0.2	0.2	2.1
G6	0.5	0.5	1.0	0.9	0.4
G7	0	0	0.1	0	0
G8	1.6	1.5	1.6	1.5	1.7



## Finding highly variable genes

scanpy

```
sc.pp.highly_variable_genes(adata, n_top_genes=2000, batch_key="sample")
```

Seurat

```
pbmc <- FindVariableFeatures(pbmc, selection.method = "vst", nfeatures = 2000)
```

Why?


	A	B	C	D	E
G4	0	0	0	0	0
G5	1.8	1.9	0.2	0.2	2.1
G6	0.5	0.5	1.0	0.9	0.4
G7	0	0	0.1	0	0
G8	1.6	1.5	1.6	1.5	1.7

- Genes that are hardly expressed at all and/or do not vary a lot across cells are less valuable for analysis
- Masking them increases computational efficiency while simultaneously reducing analysis noise



## Finding highly variable genes

How?




**Highly expressed gene**

Mean = 11.4 count

Variance = 6.64 counts


Naïve idea: Let's just take the genes with the highest variance across cells.





How?

## Finding highly variable genes




**Highly expressed gene**

Mean = 11.4 count


Variance = 6.64 counts

Naïve idea: Let's just take the genes with the highest variance across cells.



Heteroscedastic!

- Genes with higher expression also have higher variance by default
- To find the interesting genes, we need to compare their variability with that of similarly expressed genes



# Finding highly variable genes



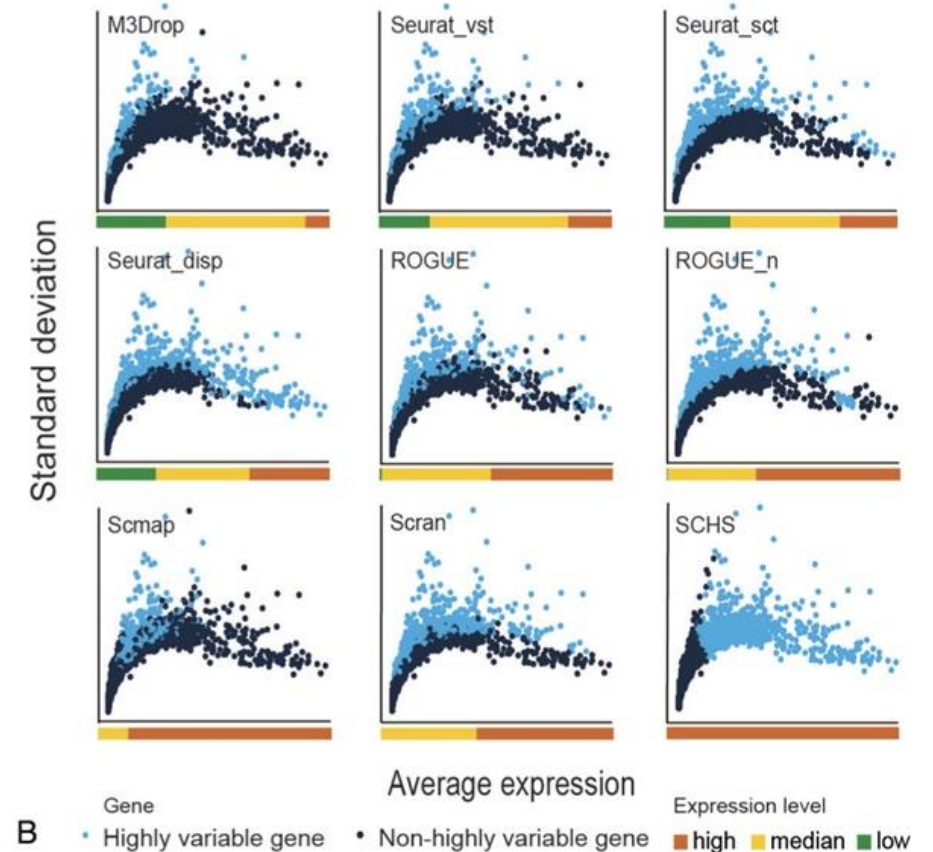
How?

typical strategy

Stabilize variance  
(e.g. log-normalize counts)

Divide genes into bins based on  
expression **or** fit a curve to the  
standard deviation over mean  
expression

Select genes which have higher  
variance than their peers



Zhang, Yanan; Xie, Xiaowei; Wu, Peng\*; Zhu, Ping\*. SIEVE: identifying robust single cell variable genes for single-cell RNA sequencing data. Blood Science 3(2):p 35-39, April 2021. | DOI: 10.1097/BS9.0000000000000072





1

2

3

## Alternative: Pearson Residuals (SCTransform)

scanpy

```
sc.experimental.pp.recipe_pearson_residuals(adata)
```

Replace normalization, log-transformation and highly variable gene search

Seurat

```
# run sctransform  
pbmc <- SCTransform(pbmc)
```

Why?

Total gene  
expression  
variability

=

Technical  
variability

+

Biological  
variability

[https://scanpy-tutorials.readthedocs.io/en/latest/tutorial\\_pearson\\_residuals.html](https://scanpy-tutorials.readthedocs.io/en/latest/tutorial_pearson_residuals.html)  
[https://satijalab.org/seurat/articles/sctransform\\_vignette](https://satijalab.org/seurat/articles/sctransform_vignette)



1

2

3

## Alternative: Pearson Residuals (SCTransform)

scanpy

```
sc.experimental.pp.recipe_pearson_residuals(adata)
```

Seurat

```
# run sctransform  
pbmc <- SCTransform(pbmc)
```

Replace normalization, log-transformation and highly variable gene search

Why?

Total gene  
expression  
variability

=

Technical  
variability

+

Biological  
variability

Biological  
variability

=

Total gene  
expression  
variability

-

Technical  
variability

measure

stats model

[https://scanpy-tutorials.readthedocs.io/en/latest/tutorial\\_pearson\\_residuals.html](https://scanpy-tutorials.readthedocs.io/en/latest/tutorial_pearson_residuals.html)  
[https://satijalab.org/seurat/articles/sctransform\\_vignette](https://satijalab.org/seurat/articles/sctransform_vignette)

1

2

3

# Alternative: Pearson Residuals (SCTransform)



How?

Biological  
variability


=

Total gene  
expression  
variability

measure

Technical  
variability

stats model



Negative binomial  
distribution describes the  
technical noise of single  
cell data.

Pearson  
residual

Raw  
count

Expected  
count under  
NB model

$$Z_{cg} = \frac{X_{cg} - \hat{\mu}_{cg}}{\sqrt{\hat{\mu}_{cg} + \hat{\mu}_{cg}^2 / \theta}}$$

Expected std  
under NB model


Lause, J., Berens, P. & Kobak, D. Analytic Pearson residuals for normalization of single-cell RNA-seq UMI data. Genome Biol 22, 258 (2021). <https://doi.org/10.1186/s13059-021-02451-7>

1

2

3

## Alternative: Pearson Residuals (SCTransform)



→ after transformation into  
Pearson Residuals

### Basic Interpretation

- **Pearson residual = 0**: The observed count matches exactly what the model expected
- **Positive residual ( $> 0$ )**: The observed count is higher than expected
- **Negative residual ( $< 0$ )**: The observed count is lower than expected

### Magnitude Interpretation

The magnitude of a Pearson residual indicates the strength of the deviation:




- **$|\text{residual}| < 2$** : Minor deviation, likely just random noise
- **$|\text{residual}| > 3$** : Strong deviation, highly likely to be biologically significant
- **$|\text{residual}| > 5$** : Extreme deviation, almost certainly represents a real biological signal

1

2

3

## Alternative: Pearson Residuals (SCTransform)






1

2

3

## Alternative: Pearson Residuals (SCTransform)



- Highly variable gene selection via Pearson Residuals can identify genes relevant for tiny cell populations
- Downstream analyses (e.g. clustering) can benefit from this
- Calculation is relatively expensive and can be prohibitive for large datasets

# Pearson Residuals are not a must



Analysis | [Open access](#) | Published: 10 April 2023

## Comparison of transformations for single-cell RNA-seq data

[Constantin Ahlmann-Eltze](#)  & [Wolfgang Huber](#)

[Nature Methods](#) **20**, 665–672 (2023) | [Cite this article](#)

**38k** Accesses | **10** Citations | **196** Altmetric | [Metrics](#)

### Abstract

The count table, a numeric matrix of genes  $\times$  cells, is the basic input data structure in the analysis of single-cell RNA-sequencing data. A common preprocessing step is to adjust the counts for variable sampling efficiency and to transform them so that the variance is similar across the dynamic range. These steps are intended to make subsequent application of generic statistical methods more palatable. Here, we describe four transformation approaches based on the **delta method, model residuals**, inferred latent expression state and factor analysis. We compare their strengths and weaknesses and find that the latter three have appealing theoretical properties; however, in benchmarks using simulated and real-world data, it turns out that a rather simple approach, namely, **the logarithm with a pseudo-count followed by principal-component analysis, performs as well or better than the more sophisticated alternatives**. This result highlights limitations of current theoretical analysis as assessed by bottom-line performance benchmarks.

Ahlmann-Eltze, C., Huber, W. Comparison of transformations for single-cell RNA-seq data. *Nat Methods* 20, 665–672 (2023).  
<https://doi.org/10.1038/s41592-023-01814-1>